



MapReduce Distributed Processing



Contents

1

Before MapReduce

2

MapReduce Overview

3

Word Count Problem

4

Word Count Flow And Solution

5

MapReduce Flow

6

Algorithm for a few Simple Problems

Before MapReduce




☁ Large scale data processing was difficult!

- Managing hundreds or thousands of processors
- Managing parallelization and distribution
- I/O Scheduling
- Status and monitoring
- Fault/crash tolerance

☁ MapReduce provides all of these, easily!

Overview of MapReduce



 Hadoop Map-Reduce is a software framework for easily writing applications which process vast amounts of data (multi-terabyte data-sets) in-parallel on large clusters (thousands of nodes) of commodity hardware in a reliable, fault-tolerant manner.

- Programming model used by Google
- A combination of the Map and Reduce models with an associated implementation
- MapReduce is highly scalable and can be used across many computers and answers all problems mentioned earlier.
- Many small machines can be used to process jobs that normally could not be processed by a large machine.
- Programming Languages Supported: Java
 - Hadoop Streaming: Python, Ruby, Perl etc
 - Hadoop Pipes: C++

MapReduce Used for



At Google:

- Index building for Google Search
- Article clustering for Google News
- Statistical machine translation

At Yahoo!:

- Index building for Yahoo! Search
- Spam detection for Yahoo! Mail

At Facebook:

- Ad optimization
 - Spam detection
-



☁ Mapper reads the data in key/value pairs

☁ Outputs zero or more key/value pairs

**map(in_key, in_value) ->
 (inter_key, inter_value) list**

☁ Mapper may use or ignore key

- The key is the byte offset into the file at which the line starts
- The value is the contents of the line itself
- Typically the key is considered irrelevant

☁ If the Mapper writes anything out, the output must be in the form of key/value pairs

Example: Turn into upper case

let map(k, v) = emit(k.toUpper(), v.toUpper())

Input - ('hadoop', 'big data') -> output - ('HADOOP', 'BIG DATA')

Reducer



- ☁ After Map phase, all intermediate values are combined together for given intermediate key into a list.
- ☁ All values associated with a particular intermediate key are guaranteed to go to the same Reducer
- ☁ The intermediate keys, and their value lists, are passed to the
- ☁ Reducer in sorted key order
- ☁ This step is known as the 'shuffle and sort'
- ☁ Reducer outputs zero or more key/value pairs

Example:

```
let reduce(k, vals)
  sum = 0
  foreach int i in vals:
    sum += i
  emit(k, sum)
```

*Input - ('bee', [111, 100, 110]) ->
Output - ('bee', 321)*

Identity Reducer:
let reduce(k, vals) =
 foreach v in vals:
 emit(k, v)

*Input - ('abc', [111, 222, 333]) -> Output
('abc', 111), ('abc', 222), ('abc', 333),*

How many Maps/Reducers?



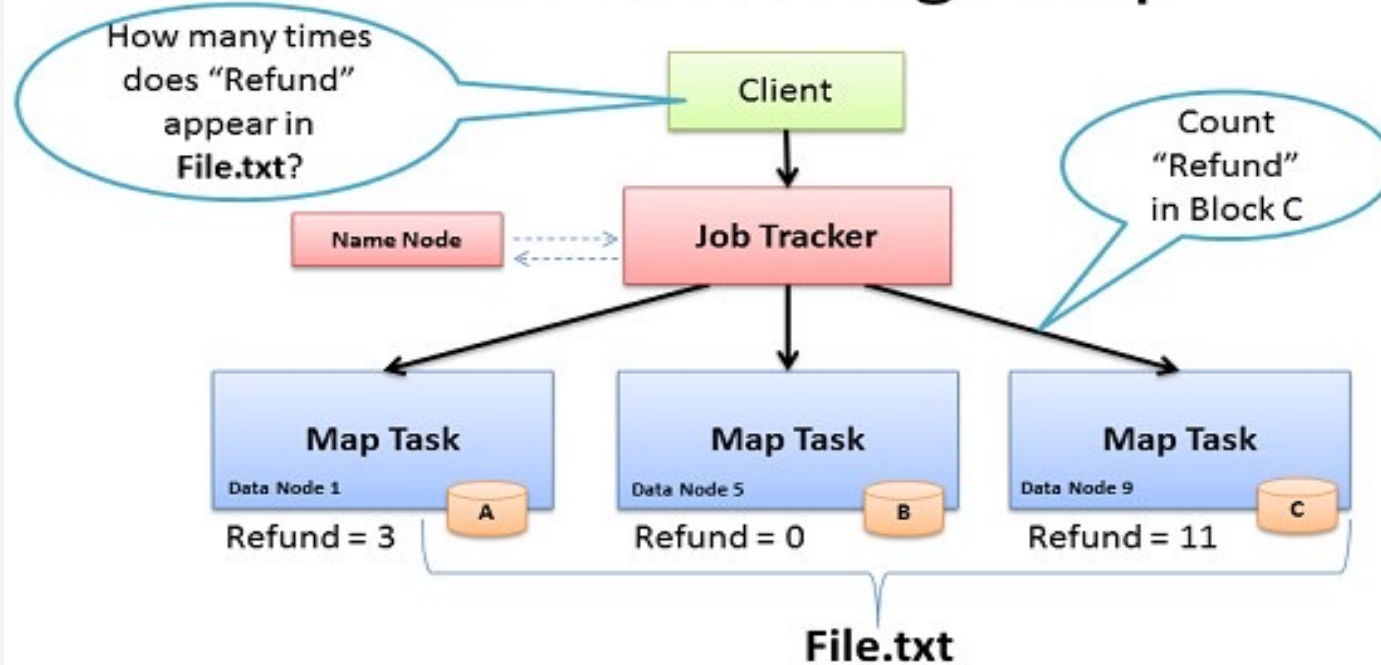
- **Maps**

- Usually as many as the number of HDFS blocks being processed, this is the default.
- Else the number of maps can be specified as a hint.
- The number of maps can also be controlled by specifying the *minimum split size*
- The actual sizes of the map inputs are computed by:
 $\text{max}(\text{min}(\text{block_size}, \text{data}/\text{\#maps}), \text{min_split_size})$

- **Reduces**

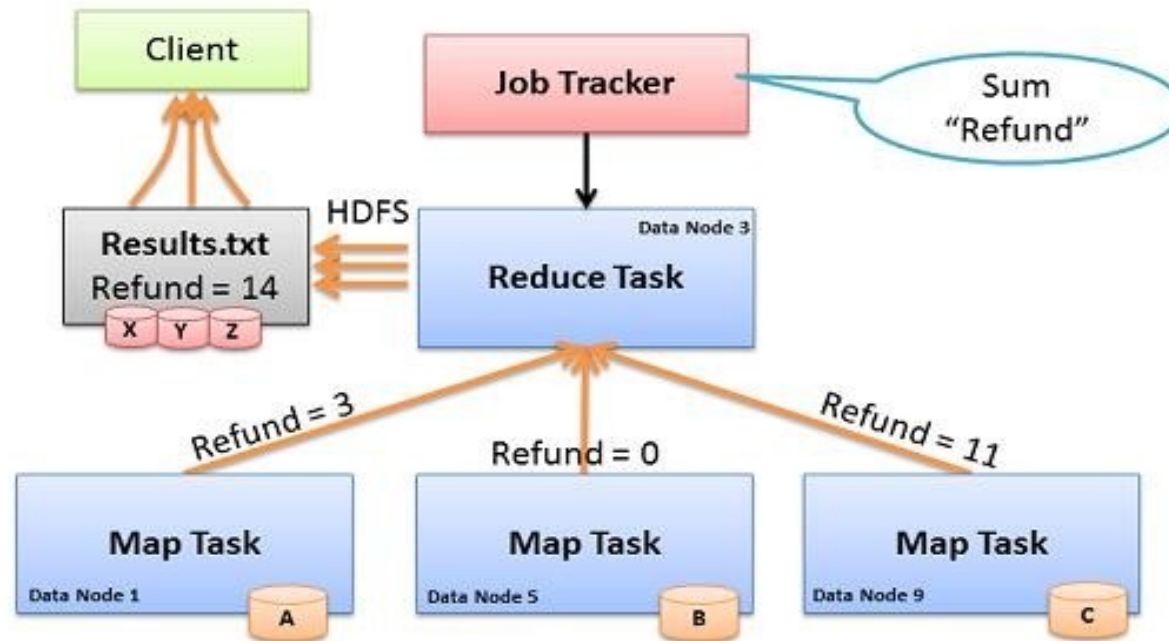
- Unless the amount of data being processed is small
 $0.95 * \text{num_nodes} * \text{mapred.tasktracker.tasks.maximum}$

Data Processing: Map



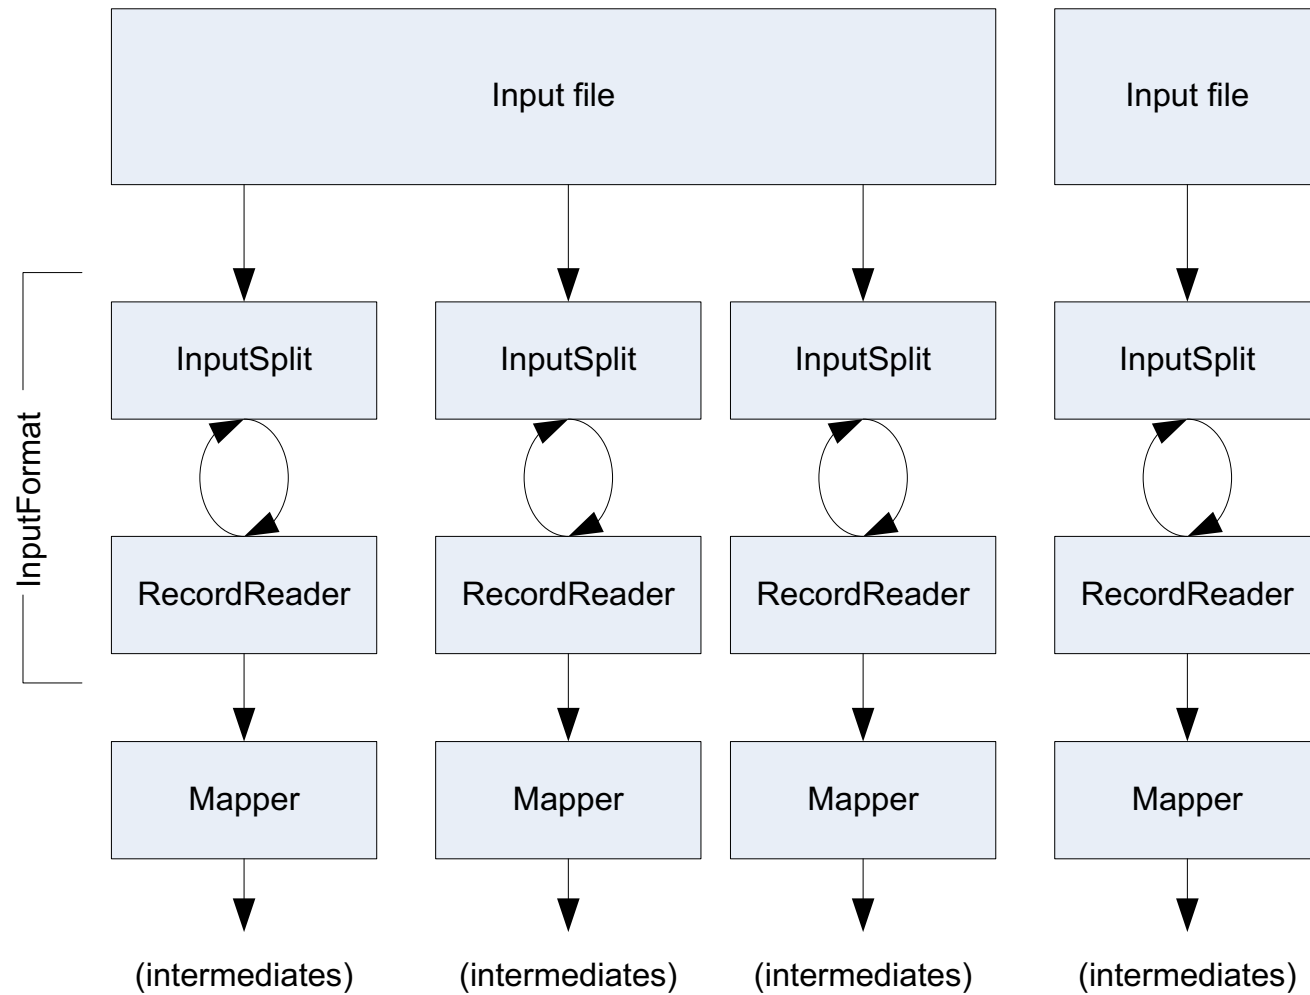
- **Map:** "Run this computation on your local data"
- Job Tracker delivers Java code to Nodes with local data

Data Processing: Reduce

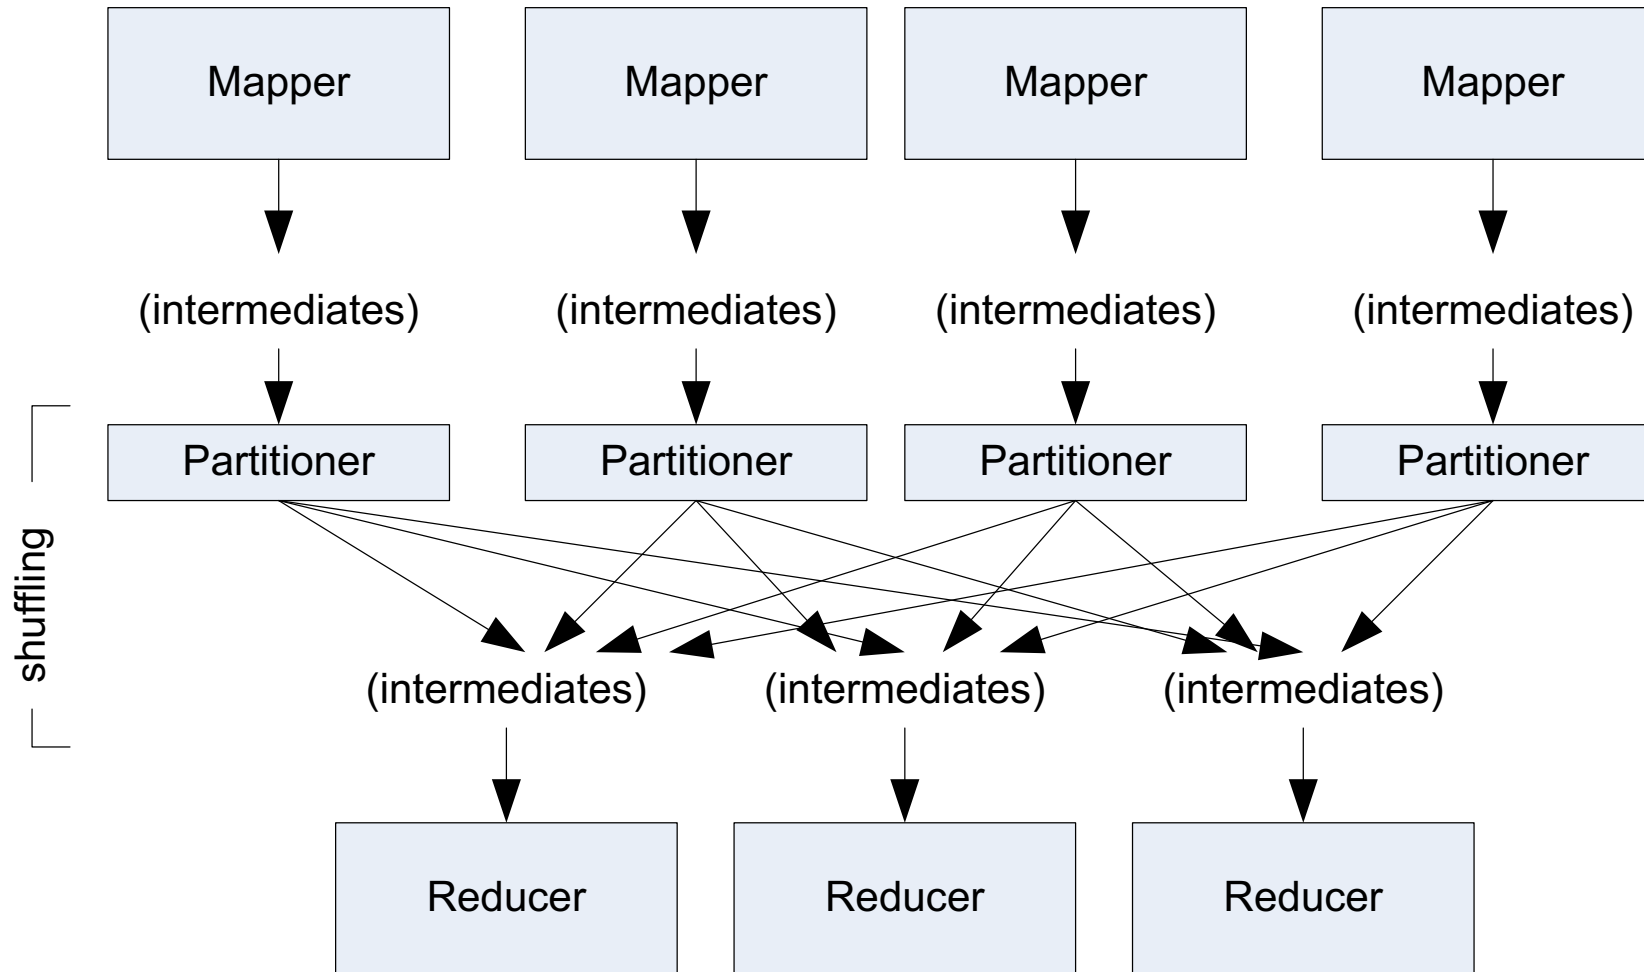


- **Reduce:** “Run this computation across Map results”
- Map Tasks send output data to Reducer over the network
- Reduce Task data output written to and read from HDFS

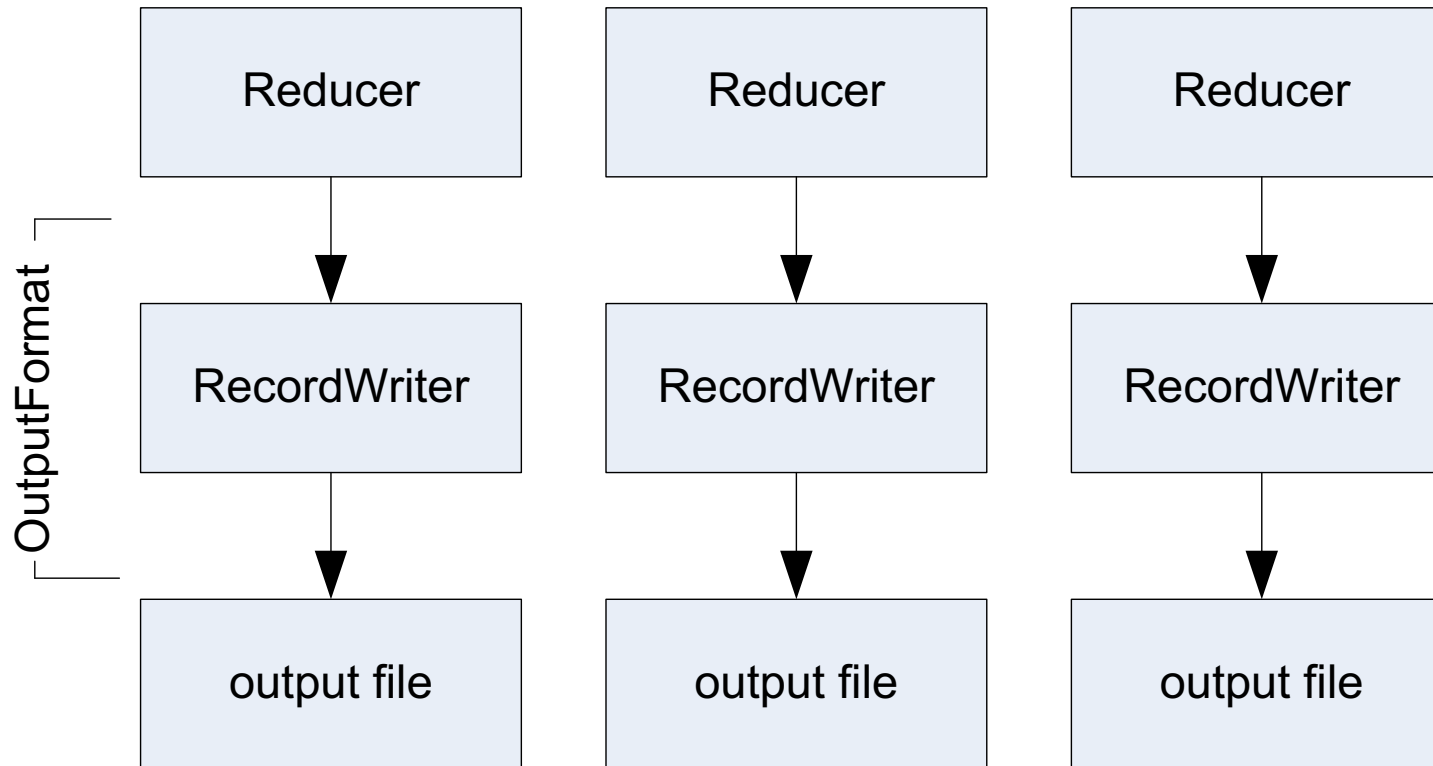
Map Reduce Flow – step1



Map Reduce Flow – Step2



Map Reduce Flow – Step3



Word Count – Distributed Solution




Problem:

Large quantity of documents

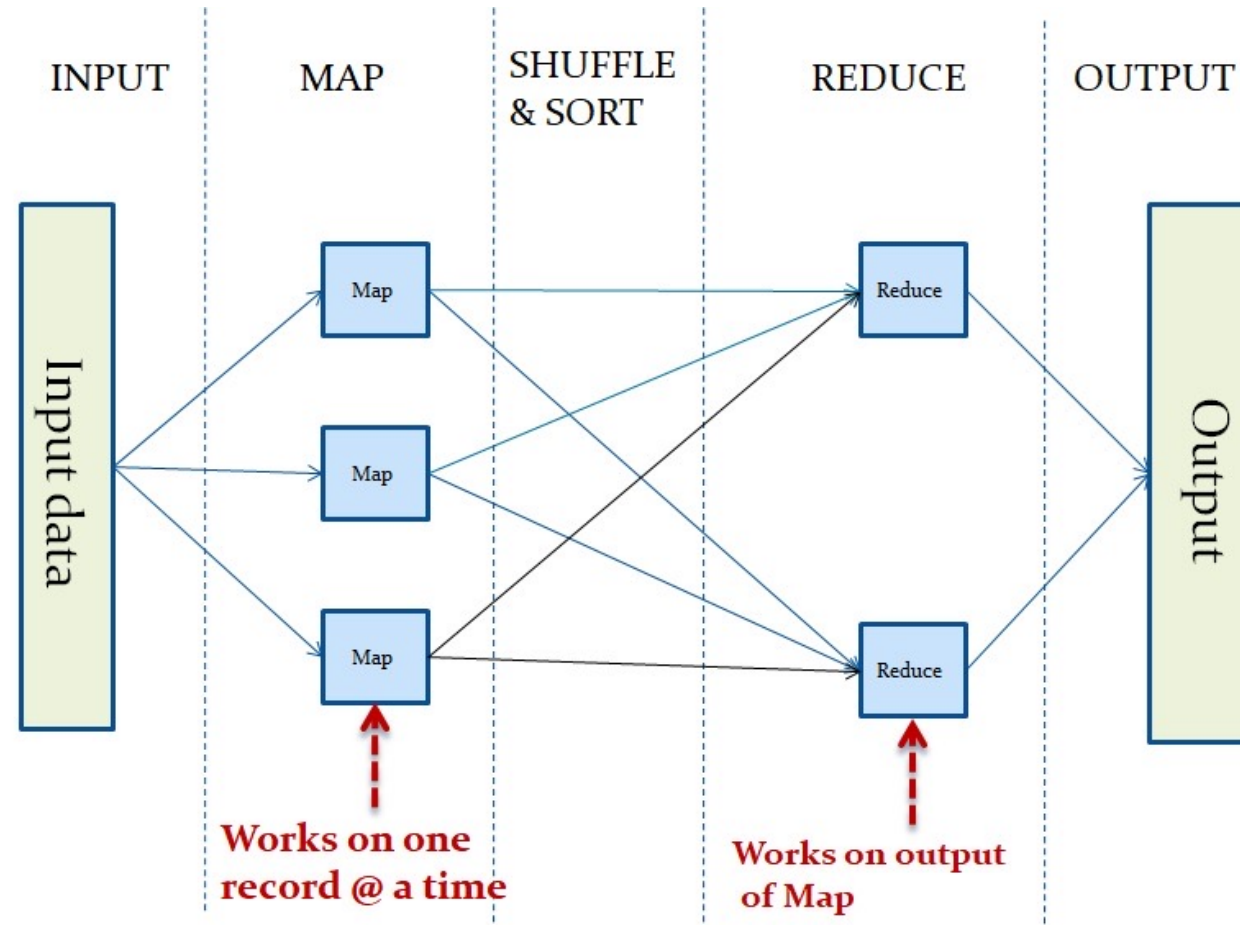
Count the number of times each distinct word occurs in the documents

Solution:

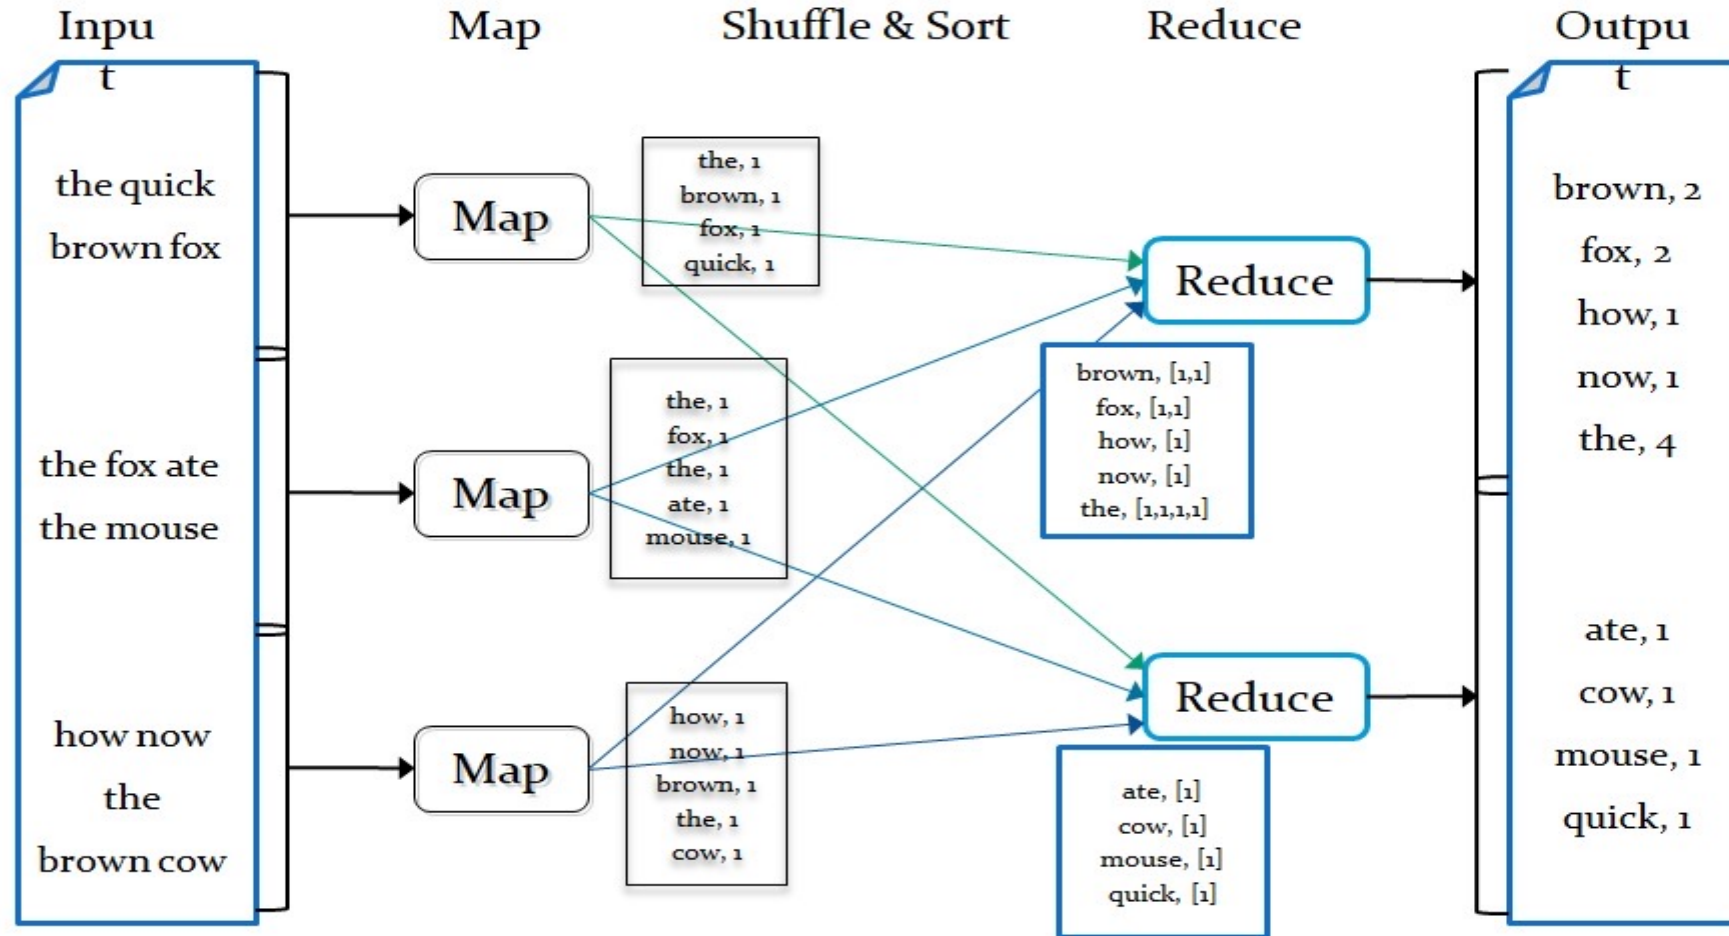
Divide and Conquer



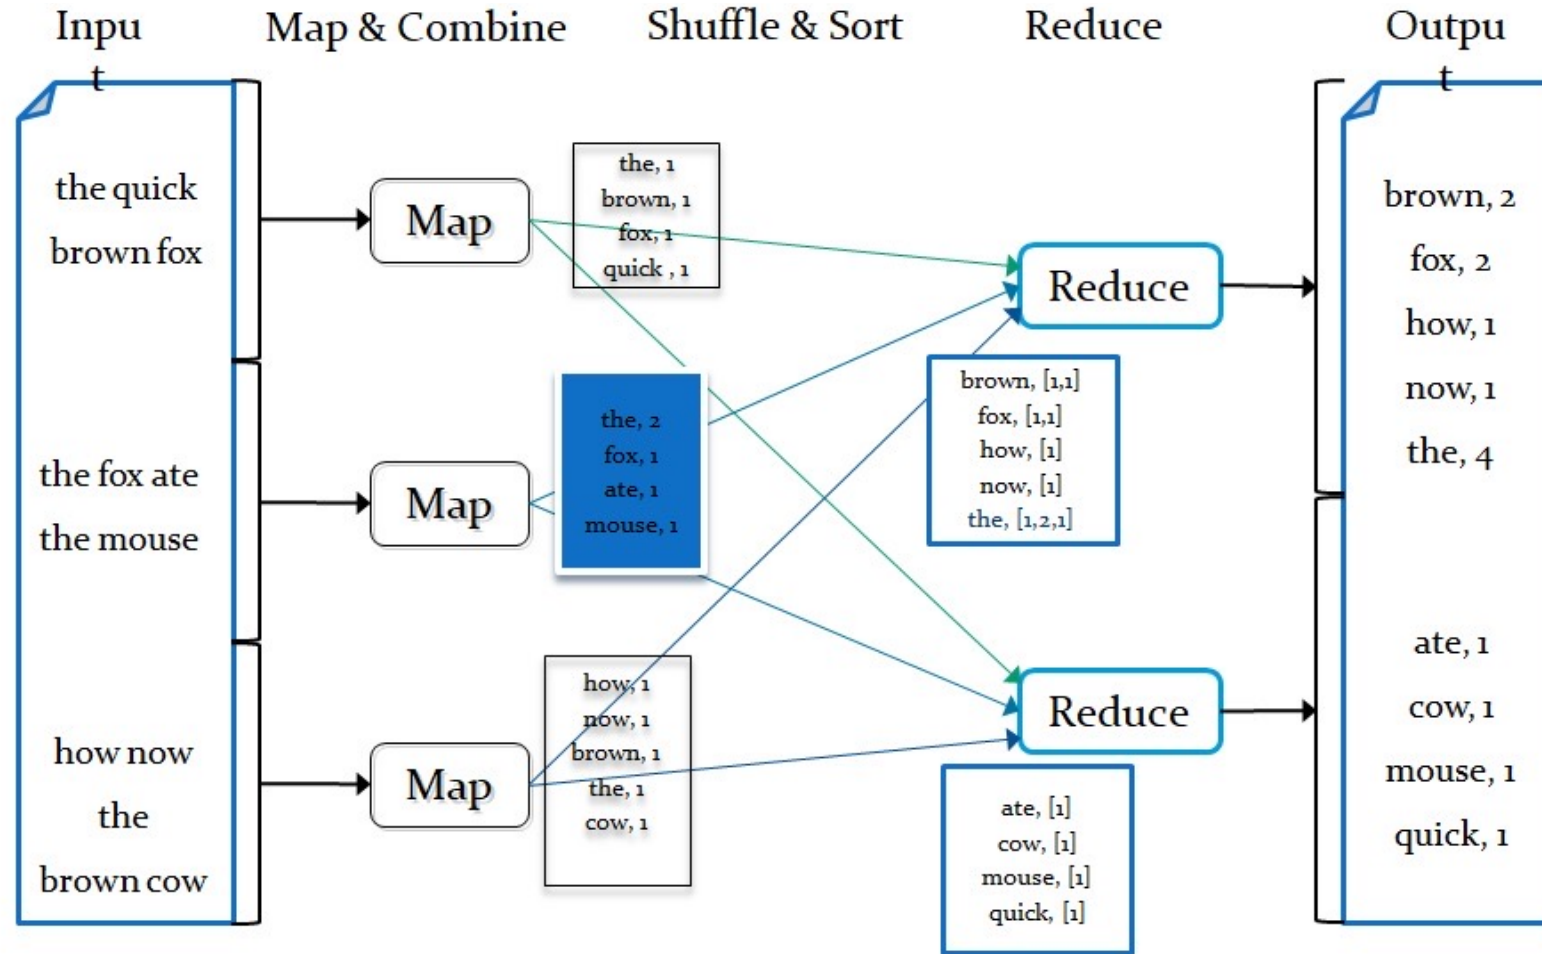
Map Reduce Flow



Word Count – Distributed Solution



Word Count – Distributed Solution





```
def mapper(line):  
    foreach word in line.split():  
        output(word, 1)
```

```
def reducer(key, values[]):  
    output(key, sum(values))
```

Map: $(k1, v1) \rightarrow \text{list } (k2, v2)$
Reduce: $(k2, \text{list } v2) \rightarrow \text{list } (k3, v3)$



Resource Requirements:

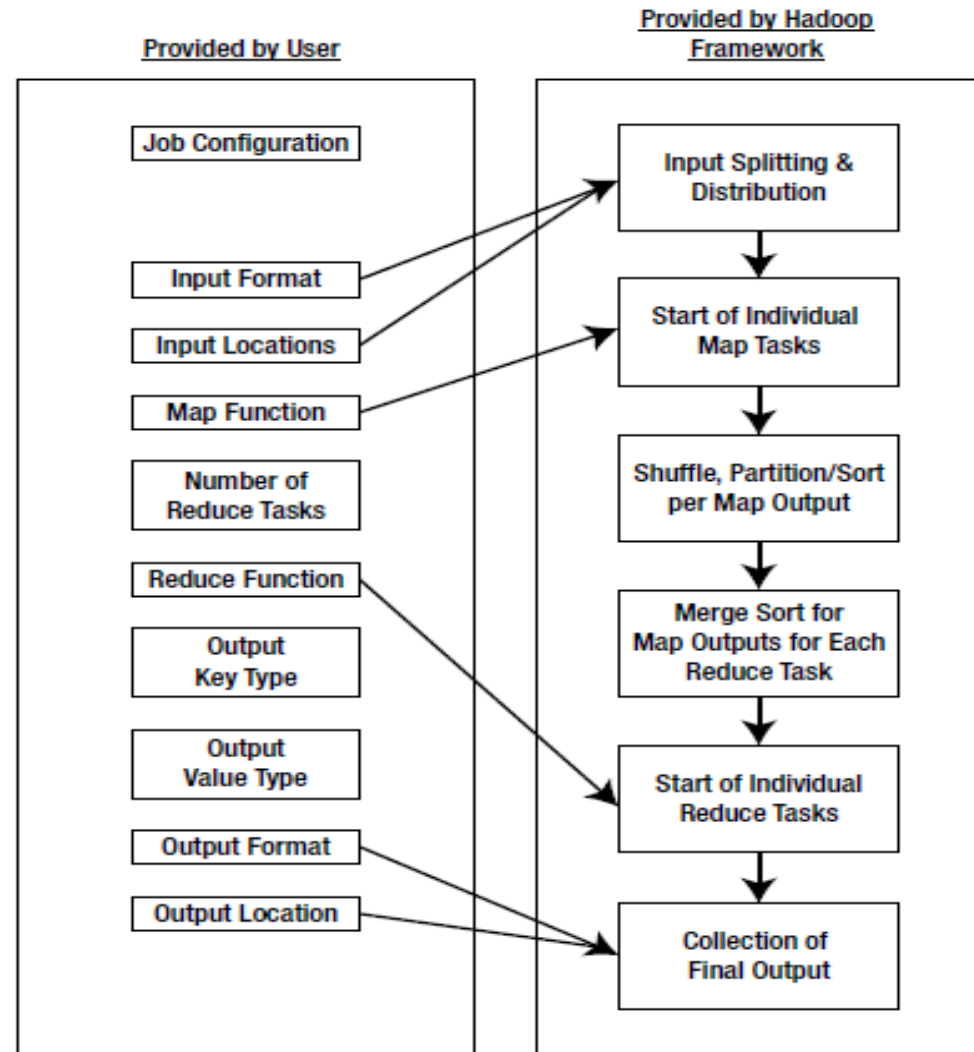
- Use of commodity machines than super computers
- Lot of storage space
- Compute near data
- Need a reliable coordinator
 - Distribute the Data
 - Coordinate the tasks

Challenges:

- A task on machine can crash, A task on machine can run very slow, A machine can crash
- There may be data loss while data being moved
- Coordinator has to be very reliable
- Corpus division amongst nodes needs to be fairly equal

WORD COUNT WAS A SIMPLE PROBLEM! IF SOMEBODY TAKES CARE OF THESE STANDARD THINGS, THEN IT CAN BE STILL KEPT SIMPLE

Parts Of MapReduce Job



Log Processing



```
.....  
10.77.25.72 [16/Feb/2009:14:29:27] http://www.google.com/  
10.77.25.72 [16/Feb/2009:14:29:27] http://www.rediff.com/  
...  
10.88.45.101 [16/Feb/2009:16:58:13] http://www.cnn.com/  
10.88.45.101 [16/Feb/2009:16:58:13] http://www.yahoo.com/  
10.77.45.82 [17/Feb/2009:09:35:42] http://news.cnet.com/  
10.88.45.101 [18/Feb/2009:16:03:07] http://www.cnn.com/  
....  
10.88.45.101 [18/Feb/2009:16:03:07] http://www.yahoo.com/  
10.88.88.69 [09/Jun/2009:15:43:50] http://www.cnn.com/  
10.88.88.69 [11/Jun/2009:16:11:25] http://news.cnet.com/  
10.88.88.69 [11/Jun/2009:16:11:25] http://www.cnn.com/  
....  
10.77.224.66 [12/Jun/2009:14:41:54] http://www.rediff.com/  
10.77.6.54 [12/Jun/2009:15:05:31] http://www.yahoo.com/  
10.88.88.69 [19/Jun/2009:16:31:11] http://www.cnn.com/  
10.77.222.22 [26/Jun/2009:15:07:46] http://www.yahoo.com/  
10.77.222.22 [26/Jun/2009:15:07:50] http://www.cnn.com/  
10.77.222.22 [26/Jun/2009:15:11:56] http://www.yahoo.com/  
10.88.88.69 [26/Jun/2009:15:48:15] http://www.cnn.com/  
10.88.88.69 [02/Jul/2009:16:11:51] http://www.yahoo.com/  
.....  
....
```

Log file may be in
Terabytes

Find Url
Frequencies

Log Processing



Key: Offset into the file
Value: 10.88.45.101 [16/Feb/2009:16:58:13]
http://www.cnn.com/



Key: http://www.cnn.com/
Value: 1



Key: http://www.cnn.com/
Value: [1,1,1,1,1/.....] or [n1, n2, n3 ...] if combiner is used



Key1: http://www.cnn.com/
Value1: 1000
Key2: http://www.yahoo.com/
Value2: 5000

Input to
Mapper

Output of
Mapper

Input to
Reducer

Output
from
Reducer



Next Session:

YARN Architecture

